



Image Denoising Using Spatial and Frequency Domain Methods

Course: Image Processing

Student Name: Joud Ali Alshehri | Student ID: 2230006037



NOISE IDENTIFICATION AND JUSTIFICATION

Based on the quantitative behavior of the filters and visual inspection of the noise degradation across the dataset, the images exhibit specific noise characteristics:

1. Camera

Likely Noise Type: Gaussian Noise (Additive White Noise).

Justification: Linear smoothing filters (like Gaussian and Mean) along with the Bilateral filter showed a steady improvement in both PSNR and SSIM as the kernel size increased, peaking at $k=7$ (PSNR = 27.75 dB, SSIM = 0.6926). Gaussian noise modifies every pixel by a small random variance; hence, local neighborhood windowing that incorporates structural weights (Bilateral) works best, whereas rank-order operators like the Median filter performed poorly here.

2. Coins

Likely Noise Type: Impulse Noise (Salt-and-Pepper Noise).

Justification: The Median filter with a 3×3 kernel achieved an exceptionally high performance (PSNR = 28.07 dB, SSIM = 0.8536), completely outclassing the Bilateral filter (SSIM = 0.3565) and the Mean filter (SSIM = 0.4905). Salt-and-pepper noise manifests as extreme black and white pixel spikes. Linear filters blur these spikes into the surrounding pixels, destroying the image, while the Median filter targets and eliminates them completely.

3. Text

Likely Noise Type: High-Density Impulse Noise (Salt-and-Pepper Noise).

Justification: The text image is dominated by severe pixel-level dropouts that corrupt the binary-like structure of text characters. The Spatial Median filter ($k=3$) achieved a massive SSIM of 0.9043 and a PSNR of 32.01 dB. Conversely, the Bilateral and Mean filters generated extremely low scores (SSIM < 0.45) because smoothing text boundaries contaminated with impulse spikes renders the characters completely unreadable.

4. Rice

Likely Noise Type: Mixed Noise / Moderate Gaussian Noise.

Justification: Unlike the text and coins images, the Gaussian filter ($k=7$) yielded the optimal response (PSNR = 24.17 dB, SSIM = 0.6214), closely followed by the Bilateral filter (SSIM = 0.5940). This indicates that the noise distribution is continuous and smoothly distributed across the soft gradients of the background and the grains, rather than being composed of isolated salt-and-pepper spikes.



5. Astronaut

Likely Noise Type: Impulse Noise (Salt-and-Pepper Noise).

Justification: The image presents severe spot-like corruptions across its colorful and highly textured regions. The Spatial Median filter ($k=3$) restored the image with near-perfect perceptual accuracy, achieving a monumental SSIM of 0.9375 and PSNR of 30.32 dB. The linear alternatives and the Bilateral filter completely failed (SSIM = 0.2682 for Bilateral), proving that rank-order statistical sorting is mandatory for this image.



METHOD SELECTION AND JUSTIFICATION

Chosen Spatial Method: Non-Linear Greedy Selection (Automated Dynamic Optimization)

To maximize the restoration quality for the evaluation, an automated optimization method was implemented in the script. It evaluates four distinct spatial filters (Median, Gaussian, Mean, and Bilateral) across three window topologies (3×3 , 5×5 , and 7×7) and automatically selects the absolute best configuration using the Structural Similarity Index (SSIM) as the primary objective function.

Engineering Rationale

Noise Model Incompatibility: No single spatial filter can effectively suppress both Gaussian and Impulse noise. Using a single fixed filter across a diverse dataset would compromise evaluation metrics.

Median Filter Mechanics: It was selected for coins, text, and astronaut because it acts as a non-linear rank-order selector. It replaces the central pixel with the median value of the window, stripping away impulse outliers without altering the surrounding clean signal.

Bilateral Filter Mechanics: It was selected for the camera image because it balances two Gaussian kernels: one in the spatial domain (closeness) and one in the radiometric/intensity domain (similarity). This allows it to smooth continuous noise while completely preserving sharp edges.

Gaussian Filter Mechanics: It was selected for the rice image because its localized distribution acts as an effective low-pass operator in the spatial domain, removing continuous variances while maintaining the macroscopic structures of the objects.



PERFORMANCE METRICS EXPLANATION (PSNR / SSIM)

To assess the restoration mathematically, two widely accepted standards in digital image processing are used:

1. Peak Signal-to-Noise Ratio (PSNR)

What it Measures: PSNR computes the ratio between the maximum possible power of the clean signal and the corrupting noise that affects its representation. It is calculated logarithmically using the Mean Squared Error (MSE) between the clean reference image and the denoised target output.

Interpretation: Higher values indicate better quality. A higher PSNR means the Mean Squared Error between the recovered image and the original image is very small, representing a lower presence of residual noise. It is expressed in decibels (dB).

2. Structural Similarity Index (SSIM)

What it Measures: Unlike PSNR, which only tracks absolute pixel-by-pixel mathematical differences, SSIM is designed to align with human visual perception. It evaluates image quality by assessing three distinct structural degradation components: luminance, contrast, and structural correlation.

Interpretation: Higher values indicate better quality. The SSIM index outputs a value bounded between -1 and +1, where a score of +1 represents a perfect structural match to the reference image. In denoising, a higher SSIM proves that the filter removed the noise without blurring out critical textures, text, or edge profiles.

RESULTS TABLE AND COMPARATIVE ANALYSIS

The table below details the quantitative performance metrics extracted from the global parameter tuning experiment.

Image Name	Evaluation Method	Optimal Tuning Parameters	PSNR (dB)	SSIM
camera	Spatial (Bilateral)	kernel_size = 7×7	27.75	0.6926
	Frequency (Gaussian LPF)	cutoff = 40	25.04	0.6879
coins	Spatial (Median)	kernel_size = 3×3	28.07	0.8536
	Frequency (Gaussian LPF)	cutoff = 40	23.38	0.5748
text	Spatial (Median)	kernel_size = 3×3	32.01	0.9043
	Frequency (Gaussian LPF)	cutoff = 20	25.59	0.6524
rice	Spatial (Gaussian)	kernel_size = 7×7	24.17	0.6214
	Frequency (Gaussian LPF)	cutoff = 40	24.14	0.6168
astronaut	Spatial (Median)	kernel_size = 3×3	30.32	0.9375
	Frequency (Gaussian LPF)	cutoff = 40	22.33	0.6044
DATASET AVG	Spatial Domain (Optimal)	Dynamic Tuning	28.46	0.8019
DATASET AVG	Frequency Domain (GLPF)	Fixed Grid	24.10	0.6273

Table 1: Detailed comparative benchmarks and dataset averages for Spatial and Frequency domain restorations.



Comprehensive Domain Comparison & Analysis

The data demonstrates that Spatial Domain Filtering completely outperforms Frequency Domain Filtering across this dataset. The spatial domain achieved a dataset average SSIM of 0.8019 and a PSNR of 28.46 dB, while the frequency domain lagged behind with an average SSIM of 0.6273 and a PSNR of 24.10 dB.

The Frequency Domain Trade-off: The Fourier-domain Gaussian Low-Pass Filter (GLPF) operates globally. In the frequency spectrum, noise and sharp structural transitions (such as text boundaries and object edges) both occupy the high-frequency range. When the GLPF cuts off high frequencies to remove noise, it unavoidably obliterates the structural edges of the image. This causes severe blurring, which drops the SSIM significantly (e.g., coins dropped from 0.8536 in spatial to 0.5748 in frequency).

The Spatial Domain Advantage: Spatial operators can be non-linear (like the Median filter). This enables them to target specific pixel corruptions based on local neighborhood statistics without altering or blurring neighboring textures. This explains why the spatial domain maintained structural preservation above 0.90 for both the text and astronaut images.

Critical Dataset Technical Note

It is structurally critical to highlight an anomaly observed in the dataset: the source directory utilized identical underlying visual assets (the *Coins* image template) for both the coins and rice testing configurations. However, despite sharing identical clean ground truths, their filtering metrics diverged dramatically during execution. The coins pipeline achieved optimal restoration via the non-linear **Median Filter**, while the rice pipeline peaked using a linear **Gaussian Filter**. This anomaly empirically validates that optimal restoration architecture is strictly dictated by the mathematical noise distribution model injected into the environment rather than the macroscopic content or high-frequency edge textures of the target image.



VISUAL RESULTS




Figure 1: Visual comparison placeholders for project benchmarking figures across the five core dataset images.

CONCLUSION

The results show that Spatial Domain Filtering Outperformed Frequency Domain Filtering in both image quality and structural preservation. The spatial methods achieved an average **PSNR of 28.46 dB** and **SSIM of 0.8019**, compared to **24.10 dB** and **0.6273** for the frequency-domain methods.

The **Median Filter** performed best for impulse noise in the *coins*, *text*, and *astronaut* images, while the **Bilateral Filter** was most effective for the *camera* image. The **Gaussian Filter** achieved the best result for the *rice* image. In contrast, the Gaussian Low-Pass Filter (GLPF) caused noticeable blurring because it removes important high-frequency image details along with noise.

Image Template	Best Spatial Method	PSNR	SSIM
camera	Bilateral	27.75	0.6926
coins	Median	28.07	0.8536
text	Median	32.01	0.9043
rice	Gaussian	24.17	0.6214
astronaut	Median	30.32	0.9375

Table 2: Summary Table

Method	Avg PSNR	Avg SSIM
Spatial Domain	28.46	0.8019
Frequency Domain	24.10	0.6273

Table 2: Dataset Average

Overall, adaptive spatial filters provided better noise removal while preserving image structures and edges. These findings highlight the importance of adaptive spatial filtering in practical image restoration applications.